



TECHNICAL LEARNING FILE

DEEP-18

Transformers, LLMs, and Retrieval

Grounded language systems

FORMAT	LENGTH	ACCESS
INTENSIVE FILE	50 PAGES	OFFLINE

Edition 2 - original curriculum - editorial review recommended



FILE / ORIENTATION

Purpose

01 FILE GUIDANCE

This optional advanced guide does not gate the core learning path. Apply transformers, llms, and retrieval with explicit assumptions, tests, tradeoffs, and limitations.

02 LOCAL USE

Index this page in the offline database and preserve its resource and page identifiers for bookmarks, notes, highlights, and progress.



FILE / ORIENTATION

Prerequisites

01 FILE GUIDANCE

Complete the related core track or pass its diagnostic.

NOUS AI ACADEMY

02 LOCAL USE

Index this page in the offline database and preserve its resource and page identifiers for bookmarks, notes, highlights, and progress.



FILE / ORIENTATION

Study Method

01 FILE GUIDANCE

For each unit, explain the concept, follow the workflow, reproduce the example, analyze a failure, and complete the applied lab. Record evidence locally.

02 LOCAL USE

Index this page in the offline database and preserve its resource and page identifiers for bookmarks, notes, highlights, and progress.



UNIT 01 / CONCEPT

Attention: Concept

01 OBJECTIVE

Explain and apply attention using explicit inputs, assumptions, steps, evidence, and failure checks.

02 CORE EXPLANATION

Attention is a central part of transformers, llms, and retrieval because it changes how inputs, assumptions, implementation choices, and evidence connect. Begin by asking what problem the idea solves, what information enters the process, what result should leave it, and what evidence would show that the result is useful. In Transformers, LLMs, and Retrieval, the terms Attention are not vocabulary to memorize in isolation; they describe a system of relationships. A strong learner can translate the idea between plain language, a diagram, a small numerical or code example, and a statement of limitations. This page establishes that translation before adding detail.

03 REASONING

Central question: When should attention be used, and what observable evidence would make that choice defensible? Build a small local example that applies attention, records the result, and tests one failure condition.

04 IMPLEMENTATION

```
for batch in train_loader:
    optimizer.zero_grad()
    prediction = model(batch.inputs)
    loss = objective(prediction, batch.targets)
    loss.backward()
    optimizer.step()
# validate on held-out data after the epoch
```

05 PRACTICE

1. Define Attention in one precise sentence and name one assumption.
2. Create a two-step example of attention and verify the intermediate state.
3. Name one boundary case, one likely misconception, and one test that would expose it.
4. Connect this page to a prior concept and explain why the connection matters.



UNIT 01 / WORKFLOW

Attention: Workflow

01 OBJECTIVE

Explain and apply attention using explicit inputs, assumptions, steps, evidence, and failure checks.

02 CORE EXPLANATION

Attention is a central part of transformers, llms, and retrieval because it changes how inputs, assumptions, implementation choices, and evidence connect. A first-principles view separates the mechanism into state, operation, constraint, and test. State describes what is currently known. The operation transforms that state. A constraint rules out invalid moves, and a test compares the output with an expectation. This structure prevents an implementation from becoming a collection of unexplained steps. It also makes debugging local: identify which state, operation, constraint, or test failed. Apply this model directly to attention.

03 REASONING

Workflow: define the smallest valid input; predict the next state; perform one operation; test one invariant; then expand. Example: Build a small local example that applies attention, records the result, and tests one failure condition.

04 IMPLEMENTATION

```
for batch in train_loader:
    optimizer.zero_grad()
    prediction = model(batch.inputs)
    loss = objective(prediction, batch.targets)
    loss.backward()
    optimizer.step()
# validate on held-out data after the epoch
```

05 PRACTICE

1. Define Attention in one precise sentence and name one assumption.
2. Create a two-step example of attention and verify the intermediate state.
3. Name one boundary case, one likely misconception, and one test that would expose it.
4. Connect this page to a prior concept and explain why the connection matters.



UNIT 01 / WORKED EXAMPLE

Attention: Worked Example

01 OBJECTIVE

Explain and apply attention using explicit inputs, assumptions, steps, evidence, and failure checks.

02 CORE EXPLANATION

Consider this task: Build a small local example that applies attention, records the result, and tests one failure condition. First state the expected output before calculating or coding. Next choose a minimal representation and execute one step at a time. After every step, compare the intermediate state with the rule that should still hold. Finally, test a boundary case that differs from the example in exactly one important way. This worked-example pattern turns attention into a repeatable method rather than a memorized answer.

03 REASONING

Worked reasoning: 1) restate the task; 2) identify knowns and unknowns; 3) choose the smallest method; 4) calculate or trace; 5) verify independently; 6) explain what would make the result fail.

04 IMPLEMENTATION

```
for batch in train_loader:
    optimizer.zero_grad()
    prediction = model(batch.inputs)
    loss = objective(prediction, batch.targets)
    loss.backward()
    optimizer.step()
# validate on held-out data after the epoch
```

05 PRACTICE

1. Define Attention in one precise sentence and name one assumption.
2. Create a two-step example of attention and verify the intermediate state.
3. Name one boundary case, one likely misconception, and one test that would expose it.
4. Connect this page to a prior concept and explain why the connection matters.



UNIT 01 / FAILURE ANALYSIS

Attention: Failure Analysis

01 OBJECTIVE

Explain and apply attention using explicit inputs, assumptions, steps, evidence, and failure checks.

02 CORE EXPLANATION

Failure analysis asks how attention can produce a plausible but wrong result. Common causes include an incorrect assumption, a representation that loses information, an edge case omitted from tests, an invalid comparison, or evidence taken from the wrong stage of the process. Change one factor at a time, preserve the failing example, and write a regression test before repairing the implementation. A clear failure record is part of the learning artifact.

03 REASONING

Failure probe: construct the smallest input that breaks the naive approach to attention. State the expected behavior and why the naive result is misleading.

04 IMPLEMENTATION

```
for batch in train_loader:
    optimizer.zero_grad()
    prediction = model(batch.inputs)
    loss = objective(prediction, batch.targets)
    loss.backward()
    optimizer.step()
# validate on held-out data after the epoch
```

05 PRACTICE

1. Define Attention in one precise sentence and name one assumption.
2. Create a two-step example of attention and verify the intermediate state.
3. Name one boundary case, one likely misconception, and one test that would expose it.
4. Connect this page to a prior concept and explain why the connection matters.



UNIT 01 / APPLIED LAB

Attention: Applied Lab

01 OBJECTIVE

Explain and apply attention using explicit inputs, assumptions, steps, evidence, and failure checks.

02 CORE EXPLANATION

Implementation converts the idea of attention into inspectable behavior. Start with a deterministic example small enough to check by hand. Keep data preparation, core logic, and reporting separate. Add assertions for shape, type, range, or state as appropriate. Record the configuration and avoid relying on hidden global state. The goal is not merely to obtain output; it is to create an output whose origin can be explained and reproduced.

03 REASONING

Lab brief: Build a small local example that applies attention, records the result, and tests one failure condition. Save the input, expected output, observed output, and a short explanation of any difference.

04 IMPLEMENTATION

```
for batch in train_loader:
    optimizer.zero_grad()
    prediction = model(batch.inputs)
    loss = objective(prediction, batch.targets)
    loss.backward()
    optimizer.step()
# validate on held-out data after the epoch
```

05 PRACTICE

1. Define Attention in one precise sentence and name one assumption.
2. Create a two-step example of attention and verify the intermediate state.
3. Name one boundary case, one likely misconception, and one test that would expose it.
4. Connect this page to a prior concept and explain why the connection matters.



UNIT 02 / CONCEPT

Transformer Blocks: Concept

01 OBJECTIVE

Explain and apply transformer blocks using explicit inputs, assumptions, steps, evidence, and failure checks.

02 CORE EXPLANATION

Transformer Blocks is a central part of transformers, llms, and retrieval because it changes how inputs, assumptions, implementation choices, and evidence connect. Begin by asking what problem the idea solves, what information enters the process, what result should leave it, and what evidence would show that the result is useful. In Transformers, LLMs, and Retrieval, the terms Transformer, Blocks are not vocabulary to memorize in isolation; they describe a system of relationships. A strong learner can translate the idea between plain language, a diagram, a small numerical or code example, and a statement of limitations. This page establishes that translation before adding detail.

03 REASONING

Central question: When should transformer blocks be used, and what observable evidence would make that choice defensible? Build a small local example that applies transformer blocks, records the result, and tests one failure condition.

04 IMPLEMENTATION

```
for batch in train_loader:
    optimizer.zero_grad()
    prediction = model(batch.inputs)
    loss = objective(prediction, batch.targets)
    loss.backward()
    optimizer.step()
# validate on held-out data after the epoch
```

05 PRACTICE

1. Define Transformer Blocks in one precise sentence and name one assumption.
2. Create a two-step example of transformer blocks and verify the intermediate state.
3. Name one boundary case, one likely misconception, and one test that would expose it.
4. Connect this page to a prior concept and explain why the connection matters.



UNIT 02 / WORKFLOW

Transformer Blocks: Workflow

01 OBJECTIVE

Explain and apply transformer blocks using explicit inputs, assumptions, steps, evidence, and failure checks.

02 CORE EXPLANATION

Transformer Blocks is a central part of transformers, llms, and retrieval because it changes how inputs, assumptions, implementation choices, and evidence connect. A first-principles view separates the mechanism into state, operation, constraint, and test. State describes what is currently known. The operation transforms that state. A constraint rules out invalid moves, and a test compares the output with an expectation. This structure prevents an implementation from becoming a collection of unexplained steps. It also makes debugging local: identify which state, operation, constraint, or test failed. Apply this model directly to transformer blocks.

03 REASONING

Workflow: define the smallest valid input; predict the next state; perform one operation; test one invariant; then expand. Example: Build a small local example that applies transformer blocks, records the result, and tests one failure condition.

04 IMPLEMENTATION

```
for batch in train_loader:
    optimizer.zero_grad()
    prediction = model(batch.inputs)
    loss = objective(prediction, batch.targets)
    loss.backward()
    optimizer.step()
# validate on held-out data after the epoch
```

05 PRACTICE

1. Define Transformer Blocks in one precise sentence and name one assumption.
2. Create a two-step example of transformer blocks and verify the intermediate state.
3. Name one boundary case, one likely misconception, and one test that would expose it.
4. Connect this page to a prior concept and explain why the connection matters.



UNIT 02 / WORKED EXAMPLE

Transformer Blocks: Worked Example

01 OBJECTIVE

Explain and apply transformer blocks using explicit inputs, assumptions, steps, evidence, and failure checks.

02 CORE EXPLANATION

Consider this task: Build a small local example that applies transformer blocks, records the result, and tests one failure condition. First state the expected output before calculating or coding. Next choose a minimal representation and execute one step at a time. After every step, compare the intermediate state with the rule that should still hold. Finally, test a boundary case that differs from the example in exactly one important way. This worked-example pattern turns transformer blocks into a repeatable method rather than a memorized answer.

03 REASONING

Worked reasoning: 1) restate the task; 2) identify knowns and unknowns; 3) choose the smallest method; 4) calculate or trace; 5) verify independently; 6) explain what would make the result fail.

04 IMPLEMENTATION

```
for batch in train_loader:
    optimizer.zero_grad()
    prediction = model(batch.inputs)
    loss = objective(prediction, batch.targets)
    loss.backward()
    optimizer.step()
# validate on held-out data after the epoch
```

05 PRACTICE

1. Define Transformer Blocks in one precise sentence and name one assumption.
2. Create a two-step example of transformer blocks and verify the intermediate state.
3. Name one boundary case, one likely misconception, and one test that would expose it.
4. Connect this page to a prior concept and explain why the connection matters.



UNIT 02 / FAILURE ANALYSIS

Transformer Blocks: Failure Analysis

01 OBJECTIVE

Explain and apply transformer blocks using explicit inputs, assumptions, steps, evidence, and failure checks.

02 CORE EXPLANATION

Failure analysis asks how transformer blocks can produce a plausible but wrong result. Common causes include an incorrect assumption, a representation that loses information, an edge case omitted from tests, an invalid comparison, or evidence taken from the wrong stage of the process. Change one factor at a time, preserve the failing example, and write a regression test before repairing the implementation. A clear failure record is part of the learning artifact.

03 REASONING

Failure probe: construct the smallest input that breaks the naive approach to transformer blocks. State the expected behavior and why the naive result is misleading.

04 IMPLEMENTATION

```
for batch in train_loader:
    optimizer.zero_grad()
    prediction = model(batch.inputs)
    loss = objective(prediction, batch.targets)
    loss.backward()
    optimizer.step()
# validate on held-out data after the epoch
```

05 PRACTICE

1. Define Transformer Blocks in one precise sentence and name one assumption.
2. Create a two-step example of transformer blocks and verify the intermediate state.
3. Name one boundary case, one likely misconception, and one test that would expose it.
4. Connect this page to a prior concept and explain why the connection matters.



UNIT 02 / APPLIED LAB

Transformer Blocks: Applied Lab

01 OBJECTIVE

Explain and apply transformer blocks using explicit inputs, assumptions, steps, evidence, and failure checks.

02 CORE EXPLANATION

Implementation converts the idea of transformer blocks into inspectable behavior. Start with a deterministic example small enough to check by hand. Keep data preparation, core logic, and reporting separate. Add assertions for shape, type, range, or state as appropriate. Record the configuration and avoid relying on hidden global state. The goal is not merely to obtain output; it is to create an output whose origin can be explained and reproduced.

03 REASONING

Lab brief: Build a small local example that applies transformer blocks, records the result, and tests one failure condition. Save the input, expected output, observed output, and a short explanation of any difference.

04 IMPLEMENTATION

```
for batch in train_loader:
    optimizer.zero_grad()
    prediction = model(batch.inputs)
    loss = objective(prediction, batch.targets)
    loss.backward()
    optimizer.step()
# validate on held-out data after the epoch
```

05 PRACTICE

1. Define Transformer Blocks in one precise sentence and name one assumption.
2. Create a two-step example of transformer blocks and verify the intermediate state.
3. Name one boundary case, one likely misconception, and one test that would expose it.
4. Connect this page to a prior concept and explain why the connection matters.



UNIT 03 / CONCEPT

Language Modeling: Concept

01 OBJECTIVE

Explain and apply language modeling using explicit inputs, assumptions, steps, evidence, and failure checks.

02 CORE EXPLANATION

Language Modeling is a central part of transformers, llms, and retrieval because it changes how inputs, assumptions, implementation choices, and evidence connect. Begin by asking what problem the idea solves, what information enters the process, what result should leave it, and what evidence would show that the result is useful. In Transformers, LLMs, and Retrieval, the terms Language, Modeling are not vocabulary to memorize in isolation; they describe a system of relationships. A strong learner can translate the idea between plain language, a diagram, a small numerical or code example, and a statement of limitations. This page establishes that translation before adding detail.

03 REASONING

Central question: When should language modeling be used, and what observable evidence would make that choice defensible? Build a small local example that applies language modeling, records the result, and tests one failure condition.

04 IMPLEMENTATION

```
for batch in train_loader:
    optimizer.zero_grad()
    prediction = model(batch.inputs)
    loss = objective(prediction, batch.targets)
    loss.backward()
    optimizer.step()
# validate on held-out data after the epoch
```

05 PRACTICE

1. Define Language Modeling in one precise sentence and name one assumption.
2. Create a two-step example of language modeling and verify the intermediate state.
3. Name one boundary case, one likely misconception, and one test that would expose it.
4. Connect this page to a prior concept and explain why the connection matters.



UNIT 03 / WORKFLOW

Language Modeling: Workflow

01 OBJECTIVE

Explain and apply language modeling using explicit inputs, assumptions, steps, evidence, and failure checks.

02 CORE EXPLANATION

Language Modeling is a central part of transformers, llms, and retrieval because it changes how inputs, assumptions, implementation choices, and evidence connect. A first-principles view separates the mechanism into state, operation, constraint, and test. State describes what is currently known. The operation transforms that state. A constraint rules out invalid moves, and a test compares the output with an expectation. This structure prevents an implementation from becoming a collection of unexplained steps. It also makes debugging local: identify which state, operation, constraint, or test failed. Apply this model directly to language modeling.

03 REASONING

Workflow: define the smallest valid input; predict the next state; perform one operation; test one invariant; then expand. Example: Build a small local example that applies language modeling, records the result, and tests one failure condition.

04 IMPLEMENTATION

```
for batch in train_loader:
    optimizer.zero_grad()
    prediction = model(batch.inputs)
    loss = objective(prediction, batch.targets)
    loss.backward()
    optimizer.step()
# validate on held-out data after the epoch
```

05 PRACTICE

1. Define Language Modeling in one precise sentence and name one assumption.
2. Create a two-step example of language modeling and verify the intermediate state.
3. Name one boundary case, one likely misconception, and one test that would expose it.
4. Connect this page to a prior concept and explain why the connection matters.



UNIT 03 / WORKED EXAMPLE

Language Modeling: Worked Example

01 OBJECTIVE

Explain and apply language modeling using explicit inputs, assumptions, steps, evidence, and failure checks.

02 CORE EXPLANATION

Consider this task: Build a small local example that applies language modeling, records the result, and tests one failure condition. First state the expected output before calculating or coding. Next choose a minimal representation and execute one step at a time. After every step, compare the intermediate state with the rule that should still hold. Finally, test a boundary case that differs from the example in exactly one important way. This worked-example pattern turns language modeling into a repeatable method rather than a memorized answer.

03 REASONING

Worked reasoning: 1) restate the task; 2) identify knowns and unknowns; 3) choose the smallest method; 4) calculate or trace; 5) verify independently; 6) explain what would make the result fail.

04 IMPLEMENTATION

```
for batch in train_loader:
    optimizer.zero_grad()
    prediction = model(batch.inputs)
    loss = objective(prediction, batch.targets)
    loss.backward()
    optimizer.step()
# validate on held-out data after the epoch
```

05 PRACTICE

1. Define Language Modeling in one precise sentence and name one assumption.
2. Create a two-step example of language modeling and verify the intermediate state.
3. Name one boundary case, one likely misconception, and one test that would expose it.
4. Connect this page to a prior concept and explain why the connection matters.



UNIT 03 / FAILURE ANALYSIS

Language Modeling: Failure Analysis

01 OBJECTIVE

Explain and apply language modeling using explicit inputs, assumptions, steps, evidence, and failure checks.

02 CORE EXPLANATION

Failure analysis asks how language modeling can produce a plausible but wrong result. Common causes include an incorrect assumption, a representation that loses information, an edge case omitted from tests, an invalid comparison, or evidence taken from the wrong stage of the process. Change one factor at a time, preserve the failing example, and write a regression test before repairing the implementation. A clear failure record is part of the learning artifact.

03 REASONING

Failure probe: construct the smallest input that breaks the naive approach to language modeling. State the expected behavior and why the naive result is misleading.

04 IMPLEMENTATION

```
for batch in train_loader:
    optimizer.zero_grad()
    prediction = model(batch.inputs)
    loss = objective(prediction, batch.targets)
    loss.backward()
    optimizer.step()
# validate on held-out data after the epoch
```

05 PRACTICE

1. Define Language Modeling in one precise sentence and name one assumption.
2. Create a two-step example of language modeling and verify the intermediate state.
3. Name one boundary case, one likely misconception, and one test that would expose it.
4. Connect this page to a prior concept and explain why the connection matters.



UNIT 03 / APPLIED LAB

Language Modeling: Applied Lab

01 OBJECTIVE

Explain and apply language modeling using explicit inputs, assumptions, steps, evidence, and failure checks.

02 CORE EXPLANATION

Implementation converts the idea of language modeling into inspectable behavior. Start with a deterministic example small enough to check by hand. Keep data preparation, core logic, and reporting separate. Add assertions for shape, type, range, or state as appropriate. Record the configuration and avoid relying on hidden global state. The goal is not merely to obtain output; it is to create an output whose origin can be explained and reproduced.

03 REASONING

Lab brief: Build a small local example that applies language modeling, records the result, and tests one failure condition. Save the input, expected output, observed output, and a short explanation of any difference.

04 IMPLEMENTATION

```
for batch in train_loader:
    optimizer.zero_grad()
    prediction = model(batch.inputs)
    loss = objective(prediction, batch.targets)
    loss.backward()
    optimizer.step()
# validate on held-out data after the epoch
```

05 PRACTICE

1. Define Language Modeling in one precise sentence and name one assumption.
2. Create a two-step example of language modeling and verify the intermediate state.
3. Name one boundary case, one likely misconception, and one test that would expose it.
4. Connect this page to a prior concept and explain why the connection matters.



UNIT 04 / CONCEPT

Embeddings: Concept

01 OBJECTIVE

Explain and apply embeddings using explicit inputs, assumptions, steps, evidence, and failure checks.

02 CORE EXPLANATION

Embeddings is a central part of transformers, llms, and retrieval because it changes how inputs, assumptions, implementation choices, and evidence connect. Begin by asking what problem the idea solves, what information enters the process, what result should leave it, and what evidence would show that the result is useful. In Transformers, LLMs, and Retrieval, the terms Embeddings are not vocabulary to memorize in isolation; they describe a system of relationships. A strong learner can translate the idea between plain language, a diagram, a small numerical or code example, and a statement of limitations. This page establishes that translation before adding detail.

03 REASONING

Central question: When should embeddings be used, and what observable evidence would make that choice defensible? Build a small local example that applies embeddings, records the result, and tests one failure condition.

04 IMPLEMENTATION

```
for batch in train_loader:
    optimizer.zero_grad()
    prediction = model(batch.inputs)
    loss = objective(prediction, batch.targets)
    loss.backward()
    optimizer.step()
# validate on held-out data after the epoch
```

05 PRACTICE

1. Define Embeddings in one precise sentence and name one assumption.
2. Create a two-step example of embeddings and verify the intermediate state.
3. Name one boundary case, one likely misconception, and one test that would expose it.
4. Connect this page to a prior concept and explain why the connection matters.



UNIT 04 / WORKFLOW

Embeddings: Workflow

01 OBJECTIVE

Explain and apply embeddings using explicit inputs, assumptions, steps, evidence, and failure checks.

02 CORE EXPLANATION

Embeddings is a central part of transformers, llms, and retrieval because it changes how inputs, assumptions, implementation choices, and evidence connect. A first-principles view separates the mechanism into state, operation, constraint, and test. State describes what is currently known. The operation transforms that state. A constraint rules out invalid moves, and a test compares the output with an expectation. This structure prevents an implementation from becoming a collection of unexplained steps. It also makes debugging local: identify which state, operation, constraint, or test failed. Apply this model directly to embeddings.

03 REASONING

Workflow: define the smallest valid input; predict the next state; perform one operation; test one invariant; then expand. Example: Build a small local example that applies embeddings, records the result, and tests one failure condition.

04 IMPLEMENTATION

```
for batch in train_loader:
    optimizer.zero_grad()
    prediction = model(batch.inputs)
    loss = objective(prediction, batch.targets)
    loss.backward()
    optimizer.step()
# validate on held-out data after the epoch
```

05 PRACTICE

1. Define Embeddings in one precise sentence and name one assumption.
2. Create a two-step example of embeddings and verify the intermediate state.
3. Name one boundary case, one likely misconception, and one test that would expose it.
4. Connect this page to a prior concept and explain why the connection matters.



UNIT 04 / WORKED EXAMPLE

Embeddings: Worked Example

01 OBJECTIVE

Explain and apply embeddings using explicit inputs, assumptions, steps, evidence, and failure checks.

02 CORE EXPLANATION

Consider this task: Build a small local example that applies embeddings, records the result, and tests one failure condition. First state the expected output before calculating or coding. Next choose a minimal representation and execute one step at a time. After every step, compare the intermediate state with the rule that should still hold. Finally, test a boundary case that differs from the example in exactly one important way. This worked-example pattern turns embeddings into a repeatable method rather than a memorized answer.

03 REASONING

Worked reasoning: 1) restate the task; 2) identify knowns and unknowns; 3) choose the smallest method; 4) calculate or trace; 5) verify independently; 6) explain what would make the result fail.

04 IMPLEMENTATION

```
for batch in train_loader:
    optimizer.zero_grad()
    prediction = model(batch.inputs)
    loss = objective(prediction, batch.targets)
    loss.backward()
    optimizer.step()
# validate on held-out data after the epoch
```

05 PRACTICE

1. Define Embeddings in one precise sentence and name one assumption.
2. Create a two-step example of embeddings and verify the intermediate state.
3. Name one boundary case, one likely misconception, and one test that would expose it.
4. Connect this page to a prior concept and explain why the connection matters.



UNIT 04 / FAILURE ANALYSIS

Embeddings: Failure Analysis

01 OBJECTIVE

Explain and apply embeddings using explicit inputs, assumptions, steps, evidence, and failure checks.

02 CORE EXPLANATION

Failure analysis asks how embeddings can produce a plausible but wrong result. Common causes include an incorrect assumption, a representation that loses information, an edge case omitted from tests, an invalid comparison, or evidence taken from the wrong stage of the process. Change one factor at a time, preserve the failing example, and write a regression test before repairing the implementation. A clear failure record is part of the learning artifact.

03 REASONING

Failure probe: construct the smallest input that breaks the naive approach to embeddings. State the expected behavior and why the naive result is misleading.

04 IMPLEMENTATION

```
for batch in train_loader:
    optimizer.zero_grad()
    prediction = model(batch.inputs)
    loss = objective(prediction, batch.targets)
    loss.backward()
    optimizer.step()
# validate on held-out data after the epoch
```

05 PRACTICE

1. Define Embeddings in one precise sentence and name one assumption.
2. Create a two-step example of embeddings and verify the intermediate state.
3. Name one boundary case, one likely misconception, and one test that would expose it.
4. Connect this page to a prior concept and explain why the connection matters.



UNIT 04 / APPLIED LAB

Embeddings: Applied Lab

01 OBJECTIVE

Explain and apply embeddings using explicit inputs, assumptions, steps, evidence, and failure checks.

02 CORE EXPLANATION

Implementation converts the idea of embeddings into inspectable behavior. Start with a deterministic example small enough to check by hand. Keep data preparation, core logic, and reporting separate. Add assertions for shape, type, range, or state as appropriate. Record the configuration and avoid relying on hidden global state. The goal is not merely to obtain output; it is to create an output whose origin can be explained and reproduced.

03 REASONING

Lab brief: Build a small local example that applies embeddings, records the result, and tests one failure condition. Save the input, expected output, observed output, and a short explanation of any difference.

04 IMPLEMENTATION

```
for batch in train_loader:
    optimizer.zero_grad()
    prediction = model(batch.inputs)
    loss = objective(prediction, batch.targets)
    loss.backward()
    optimizer.step()
# validate on held-out data after the epoch
```

05 PRACTICE

1. Define Embeddings in one precise sentence and name one assumption.
2. Create a two-step example of embeddings and verify the intermediate state.
3. Name one boundary case, one likely misconception, and one test that would expose it.
4. Connect this page to a prior concept and explain why the connection matters.



UNIT 05 / CONCEPT

Semantic Search: Concept

01 OBJECTIVE

Explain and apply semantic search using explicit inputs, assumptions, steps, evidence, and failure checks.

02 CORE EXPLANATION

Semantic Search is a central part of transformers, llms, and retrieval because it changes how inputs, assumptions, implementation choices, and evidence connect. Begin by asking what problem the idea solves, what information enters the process, what result should leave it, and what evidence would show that the result is useful. In Transformers, LLMs, and Retrieval, the terms Semantic, Search are not vocabulary to memorize in isolation; they describe a system of relationships. A strong learner can translate the idea between plain language, a diagram, a small numerical or code example, and a statement of limitations. This page establishes that translation before adding detail.

03 REASONING

Central question: When should semantic search be used, and what observable evidence would make that choice defensible? Build a small local example that applies semantic search, records the result, and tests one failure condition.

04 IMPLEMENTATION

```
for batch in train_loader:
    optimizer.zero_grad()
    prediction = model(batch.inputs)
    loss = objective(prediction, batch.targets)
    loss.backward()
    optimizer.step()
# validate on held-out data after the epoch
```

05 PRACTICE

1. Define Semantic Search in one precise sentence and name one assumption.
2. Create a two-step example of semantic search and verify the intermediate state.
3. Name one boundary case, one likely misconception, and one test that would expose it.
4. Connect this page to a prior concept and explain why the connection matters.



UNIT 05 / WORKFLOW

Semantic Search: Workflow

01 OBJECTIVE

Explain and apply semantic search using explicit inputs, assumptions, steps, evidence, and failure checks.

02 CORE EXPLANATION

Semantic Search is a central part of transformers, llms, and retrieval because it changes how inputs, assumptions, implementation choices, and evidence connect. A first-principles view separates the mechanism into state, operation, constraint, and test. State describes what is currently known. The operation transforms that state. A constraint rules out invalid moves, and a test compares the output with an expectation. This structure prevents an implementation from becoming a collection of unexplained steps. It also makes debugging local: identify which state, operation, constraint, or test failed. Apply this model directly to semantic search.

03 REASONING

Workflow: define the smallest valid input; predict the next state; perform one operation; test one invariant; then expand. Example: Build a small local example that applies semantic search, records the result, and tests one failure condition.

04 IMPLEMENTATION

```
for batch in train_loader:
    optimizer.zero_grad()
    prediction = model(batch.inputs)
    loss = objective(prediction, batch.targets)
    loss.backward()
    optimizer.step()
# validate on held-out data after the epoch
```

05 PRACTICE

1. Define Semantic Search in one precise sentence and name one assumption.
2. Create a two-step example of semantic search and verify the intermediate state.
3. Name one boundary case, one likely misconception, and one test that would expose it.
4. Connect this page to a prior concept and explain why the connection matters.



UNIT 05 / WORKED EXAMPLE

Semantic Search: Worked Example

01 OBJECTIVE

Explain and apply semantic search using explicit inputs, assumptions, steps, evidence, and failure checks.

02 CORE EXPLANATION

Consider this task: Build a small local example that applies semantic search, records the result, and tests one failure condition. First state the expected output before calculating or coding. Next choose a minimal representation and execute one step at a time. After every step, compare the intermediate state with the rule that should still hold. Finally, test a boundary case that differs from the example in exactly one important way. This worked-example pattern turns semantic search into a repeatable method rather than a memorized answer.

03 REASONING

Worked reasoning: 1) restate the task; 2) identify knowns and unknowns; 3) choose the smallest method; 4) calculate or trace; 5) verify independently; 6) explain what would make the result fail.

04 IMPLEMENTATION

```
for batch in train_loader:
    optimizer.zero_grad()
    prediction = model(batch.inputs)
    loss = objective(prediction, batch.targets)
    loss.backward()
    optimizer.step()
# validate on held-out data after the epoch
```

05 PRACTICE

1. Define Semantic Search in one precise sentence and name one assumption.
2. Create a two-step example of semantic search and verify the intermediate state.
3. Name one boundary case, one likely misconception, and one test that would expose it.
4. Connect this page to a prior concept and explain why the connection matters.



UNIT 05 / FAILURE ANALYSIS

Semantic Search: Failure Analysis

01 OBJECTIVE

Explain and apply semantic search using explicit inputs, assumptions, steps, evidence, and failure checks.

02 CORE EXPLANATION

Failure analysis asks how semantic search can produce a plausible but wrong result. Common causes include an incorrect assumption, a representation that loses information, an edge case omitted from tests, an invalid comparison, or evidence taken from the wrong stage of the process. Change one factor at a time, preserve the failing example, and write a regression test before repairing the implementation. A clear failure record is part of the learning artifact.

03 REASONING

Failure probe: construct the smallest input that breaks the naive approach to semantic search. State the expected behavior and why the naive result is misleading.

04 IMPLEMENTATION

```
for batch in train_loader:
    optimizer.zero_grad()
    prediction = model(batch.inputs)
    loss = objective(prediction, batch.targets)
    loss.backward()
    optimizer.step()
# validate on held-out data after the epoch
```

05 PRACTICE

1. Define Semantic Search in one precise sentence and name one assumption.
2. Create a two-step example of semantic search and verify the intermediate state.
3. Name one boundary case, one likely misconception, and one test that would expose it.
4. Connect this page to a prior concept and explain why the connection matters.



UNIT 05 / APPLIED LAB

Semantic Search: Applied Lab

01 OBJECTIVE

Explain and apply semantic search using explicit inputs, assumptions, steps, evidence, and failure checks.

02 CORE EXPLANATION

Implementation converts the idea of semantic search into inspectable behavior. Start with a deterministic example small enough to check by hand. Keep data preparation, core logic, and reporting separate. Add assertions for shape, type, range, or state as appropriate. Record the configuration and avoid relying on hidden global state. The goal is not merely to obtain output; it is to create an output whose origin can be explained and reproduced.

03 REASONING

Lab brief: Build a small local example that applies semantic search, records the result, and tests one failure condition. Save the input, expected output, observed output, and a short explanation of any difference.

04 IMPLEMENTATION

```
for batch in train_loader:
    optimizer.zero_grad()
    prediction = model(batch.inputs)
    loss = objective(prediction, batch.targets)
    loss.backward()
    optimizer.step()
# validate on held-out data after the epoch
```

05 PRACTICE

1. Define Semantic Search in one precise sentence and name one assumption.
2. Create a two-step example of semantic search and verify the intermediate state.
3. Name one boundary case, one likely misconception, and one test that would expose it.
4. Connect this page to a prior concept and explain why the connection matters.



UNIT 06 / CONCEPT

Retrieval Pipelines: Concept

01 OBJECTIVE

Explain and apply retrieval pipelines using explicit inputs, assumptions, steps, evidence, and failure checks.

02 CORE EXPLANATION

Retrieval Pipelines is a central part of transformers, llms, and retrieval because it changes how inputs, assumptions, implementation choices, and evidence connect. Begin by asking what problem the idea solves, what information enters the process, what result should leave it, and what evidence would show that the result is useful. In Transformers, LLMs, and Retrieval, the terms Retrieval, Pipelines are not vocabulary to memorize in isolation; they describe a system of relationships. A strong learner can translate the idea between plain language, a diagram, a small numerical or code example, and a statement of limitations. This page establishes that translation before adding detail.

03 REASONING

Central question: When should retrieval pipelines be used, and what observable evidence would make that choice defensible? Build a small local example that applies retrieval pipelines, records the result, and tests one failure condition.

04 IMPLEMENTATION

```
for batch in train_loader:
    optimizer.zero_grad()
    prediction = model(batch.inputs)
    loss = objective(prediction, batch.targets)
    loss.backward()
    optimizer.step()
# validate on held-out data after the epoch
```

05 PRACTICE

1. Define Retrieval Pipelines in one precise sentence and name one assumption.
2. Create a two-step example of retrieval pipelines and verify the intermediate state.
3. Name one boundary case, one likely misconception, and one test that would expose it.
4. Connect this page to a prior concept and explain why the connection matters.



UNIT 06 / WORKFLOW

Retrieval Pipelines: Workflow

01 OBJECTIVE

Explain and apply retrieval pipelines using explicit inputs, assumptions, steps, evidence, and failure checks.

02 CORE EXPLANATION

Retrieval Pipelines is a central part of transformers, llms, and retrieval because it changes how inputs, assumptions, implementation choices, and evidence connect. A first-principles view separates the mechanism into state, operation, constraint, and test. State describes what is currently known. The operation transforms that state. A constraint rules out invalid moves, and a test compares the output with an expectation. This structure prevents an implementation from becoming a collection of unexplained steps. It also makes debugging local: identify which state, operation, constraint, or test failed. Apply this model directly to retrieval pipelines.

03 REASONING

Workflow: define the smallest valid input; predict the next state; perform one operation; test one invariant; then expand. Example: Build a small local example that applies retrieval pipelines, records the result, and tests one failure condition.

04 IMPLEMENTATION

```
for batch in train_loader:
    optimizer.zero_grad()
    prediction = model(batch.inputs)
    loss = objective(prediction, batch.targets)
    loss.backward()
    optimizer.step()
# validate on held-out data after the epoch
```

05 PRACTICE

1. Define Retrieval Pipelines in one precise sentence and name one assumption.
2. Create a two-step example of retrieval pipelines and verify the intermediate state.
3. Name one boundary case, one likely misconception, and one test that would expose it.
4. Connect this page to a prior concept and explain why the connection matters.



UNIT 06 / WORKED EXAMPLE

Retrieval Pipelines: Worked Example

01 OBJECTIVE

Explain and apply retrieval pipelines using explicit inputs, assumptions, steps, evidence, and failure checks.

02 CORE EXPLANATION

Consider this task: Build a small local example that applies retrieval pipelines, records the result, and tests one failure condition. First state the expected output before calculating or coding. Next choose a minimal representation and execute one step at a time. After every step, compare the intermediate state with the rule that should still hold. Finally, test a boundary case that differs from the example in exactly one important way. This worked-example pattern turns retrieval pipelines into a repeatable method rather than a memorized answer.

03 REASONING

Worked reasoning: 1) restate the task; 2) identify knowns and unknowns; 3) choose the smallest method; 4) calculate or trace; 5) verify independently; 6) explain what would make the result fail.

04 IMPLEMENTATION

```
for batch in train_loader:
    optimizer.zero_grad()
    prediction = model(batch.inputs)
    loss = objective(prediction, batch.targets)
    loss.backward()
    optimizer.step()
# validate on held-out data after the epoch
```

05 PRACTICE

1. Define Retrieval Pipelines in one precise sentence and name one assumption.
2. Create a two-step example of retrieval pipelines and verify the intermediate state.
3. Name one boundary case, one likely misconception, and one test that would expose it.
4. Connect this page to a prior concept and explain why the connection matters.



UNIT 06 / FAILURE ANALYSIS

Retrieval Pipelines: Failure Analysis

01 OBJECTIVE

Explain and apply retrieval pipelines using explicit inputs, assumptions, steps, evidence, and failure checks.

02 CORE EXPLANATION

Failure analysis asks how retrieval pipelines can produce a plausible but wrong result. Common causes include an incorrect assumption, a representation that loses information, an edge case omitted from tests, an invalid comparison, or evidence taken from the wrong stage of the process. Change one factor at a time, preserve the failing example, and write a regression test before repairing the implementation. A clear failure record is part of the learning artifact.

03 REASONING

Failure probe: construct the smallest input that breaks the naive approach to retrieval pipelines. State the expected behavior and why the naive result is misleading.

04 IMPLEMENTATION

```
for batch in train_loader:
    optimizer.zero_grad()
    prediction = model(batch.inputs)
    loss = objective(prediction, batch.targets)
    loss.backward()
    optimizer.step()
# validate on held-out data after the epoch
```

05 PRACTICE

1. Define Retrieval Pipelines in one precise sentence and name one assumption.
2. Create a two-step example of retrieval pipelines and verify the intermediate state.
3. Name one boundary case, one likely misconception, and one test that would expose it.
4. Connect this page to a prior concept and explain why the connection matters.



UNIT 06 / APPLIED LAB

Retrieval Pipelines: Applied Lab

01 OBJECTIVE

Explain and apply retrieval pipelines using explicit inputs, assumptions, steps, evidence, and failure checks.

02 CORE EXPLANATION

Implementation converts the idea of retrieval pipelines into inspectable behavior. Start with a deterministic example small enough to check by hand. Keep data preparation, core logic, and reporting separate. Add assertions for shape, type, range, or state as appropriate. Record the configuration and avoid relying on hidden global state. The goal is not merely to obtain output; it is to create an output whose origin can be explained and reproduced.

03 REASONING

Lab brief: Build a small local example that applies retrieval pipelines, records the result, and tests one failure condition. Save the input, expected output, observed output, and a short explanation of any difference.

04 IMPLEMENTATION

```
for batch in train_loader:
    optimizer.zero_grad()
    prediction = model(batch.inputs)
    loss = objective(prediction, batch.targets)
    loss.backward()
    optimizer.step()
# validate on held-out data after the epoch
```

05 PRACTICE

1. Define Retrieval Pipelines in one precise sentence and name one assumption.
2. Create a two-step example of retrieval pipelines and verify the intermediate state.
3. Name one boundary case, one likely misconception, and one test that would expose it.
4. Connect this page to a prior concept and explain why the connection matters.



UNIT 07 / CONCEPT

Tool Use: Concept

01 OBJECTIVE

Explain and apply tool use using explicit inputs, assumptions, steps, evidence, and failure checks.

02 CORE EXPLANATION

Tool Use is a central part of transformers, llms, and retrieval because it changes how inputs, assumptions, implementation choices, and evidence connect. Begin by asking what problem the idea solves, what information enters the process, what result should leave it, and what evidence would show that the result is useful. In Transformers, LLMs, and Retrieval, the terms Tool are not vocabulary to memorize in isolation; they describe a system of relationships. A strong learner can translate the idea between plain language, a diagram, a small numerical or code example, and a statement of limitations. This page establishes that translation before adding detail.

03 REASONING

Central question: When should tool use be used, and what observable evidence would make that choice defensible? Build a small local example that applies tool use, records the result, and tests one failure condition.

04 IMPLEMENTATION

```
for batch in train_loader:
    optimizer.zero_grad()
    prediction = model(batch.inputs)
    loss = objective(prediction, batch.targets)
    loss.backward()
    optimizer.step()
# validate on held-out data after the epoch
```

05 PRACTICE

1. Define Tool Use in one precise sentence and name one assumption.
2. Create a two-step example of tool use and verify the intermediate state.
3. Name one boundary case, one likely misconception, and one test that would expose it.
4. Connect this page to a prior concept and explain why the connection matters.



UNIT 07 / WORKFLOW

Tool Use: Workflow

01 OBJECTIVE

Explain and apply tool use using explicit inputs, assumptions, steps, evidence, and failure checks.

02 CORE EXPLANATION

Tool Use is a central part of transformers, llms, and retrieval because it changes how inputs, assumptions, implementation choices, and evidence connect. A first-principles view separates the mechanism into state, operation, constraint, and test. State describes what is currently known. The operation transforms that state. A constraint rules out invalid moves, and a test compares the output with an expectation. This structure prevents an implementation from becoming a collection of unexplained steps. It also makes debugging local: identify which state, operation, constraint, or test failed. Apply this model directly to tool use.

03 REASONING

Workflow: define the smallest valid input; predict the next state; perform one operation; test one invariant; then expand. Example: Build a small local example that applies tool use, records the result, and tests one failure condition.

04 IMPLEMENTATION

```
for batch in train_loader:
    optimizer.zero_grad()
    prediction = model(batch.inputs)
    loss = objective(prediction, batch.targets)
    loss.backward()
    optimizer.step()
# validate on held-out data after the epoch
```

05 PRACTICE

1. Define Tool Use in one precise sentence and name one assumption.
2. Create a two-step example of tool use and verify the intermediate state.
3. Name one boundary case, one likely misconception, and one test that would expose it.
4. Connect this page to a prior concept and explain why the connection matters.



UNIT 07 / WORKED EXAMPLE

Tool Use: Worked Example

01 OBJECTIVE

Explain and apply tool use using explicit inputs, assumptions, steps, evidence, and failure checks.

02 CORE EXPLANATION

Consider this task: Build a small local example that applies tool use, records the result, and tests one failure condition. First state the expected output before calculating or coding. Next choose a minimal representation and execute one step at a time. After every step, compare the intermediate state with the rule that should still hold. Finally, test a boundary case that differs from the example in exactly one important way. This worked-example pattern turns tool use into a repeatable method rather than a memorized answer.

03 REASONING

Worked reasoning: 1) restate the task; 2) identify knowns and unknowns; 3) choose the smallest method; 4) calculate or trace; 5) verify independently; 6) explain what would make the result fail.

04 IMPLEMENTATION

```
for batch in train_loader:
    optimizer.zero_grad()
    prediction = model(batch.inputs)
    loss = objective(prediction, batch.targets)
    loss.backward()
    optimizer.step()
# validate on held-out data after the epoch
```

05 PRACTICE

1. Define Tool Use in one precise sentence and name one assumption.
2. Create a two-step example of tool use and verify the intermediate state.
3. Name one boundary case, one likely misconception, and one test that would expose it.
4. Connect this page to a prior concept and explain why the connection matters.



UNIT 07 / FAILURE ANALYSIS

Tool Use: Failure Analysis

01 OBJECTIVE

Explain and apply tool use using explicit inputs, assumptions, steps, evidence, and failure checks.

02 CORE EXPLANATION

Failure analysis asks how tool use can produce a plausible but wrong result. Common causes include an incorrect assumption, a representation that loses information, an edge case omitted from tests, an invalid comparison, or evidence taken from the wrong stage of the process. Change one factor at a time, preserve the failing example, and write a regression test before repairing the implementation. A clear failure record is part of the learning artifact.

03 REASONING

Failure probe: construct the smallest input that breaks the naive approach to tool use. State the expected behavior and why the naive result is misleading.

04 IMPLEMENTATION

```
for batch in train_loader:
    optimizer.zero_grad()
    prediction = model(batch.inputs)
    loss = objective(prediction, batch.targets)
    loss.backward()
    optimizer.step()
# validate on held-out data after the epoch
```

05 PRACTICE

1. Define Tool Use in one precise sentence and name one assumption.
2. Create a two-step example of tool use and verify the intermediate state.
3. Name one boundary case, one likely misconception, and one test that would expose it.
4. Connect this page to a prior concept and explain why the connection matters.



UNIT 07 / APPLIED LAB

Tool Use: Applied Lab

01 OBJECTIVE

Explain and apply tool use using explicit inputs, assumptions, steps, evidence, and failure checks.

02 CORE EXPLANATION

Implementation converts the idea of tool use into inspectable behavior. Start with a deterministic example small enough to check by hand. Keep data preparation, core logic, and reporting separate. Add assertions for shape, type, range, or state as appropriate. Record the configuration and avoid relying on hidden global state. The goal is not merely to obtain output; it is to create an output whose origin can be explained and reproduced.

03 REASONING

Lab brief: Build a small local example that applies tool use, records the result, and tests one failure condition. Save the input, expected output, observed output, and a short explanation of any difference.

04 IMPLEMENTATION

```
for batch in train_loader:
    optimizer.zero_grad()
    prediction = model(batch.inputs)
    loss = objective(prediction, batch.targets)
    loss.backward()
    optimizer.step()
# validate on held-out data after the epoch
```

05 PRACTICE

1. Define Tool Use in one precise sentence and name one assumption.
2. Create a two-step example of tool use and verify the intermediate state.
3. Name one boundary case, one likely misconception, and one test that would expose it.
4. Connect this page to a prior concept and explain why the connection matters.



UNIT 08 / CONCEPT

Evaluation: Concept

01 OBJECTIVE

Explain and apply evaluation using explicit inputs, assumptions, steps, evidence, and failure checks.

02 CORE EXPLANATION

Evaluation is a central part of transformers, llms, and retrieval because it changes how inputs, assumptions, implementation choices, and evidence connect. Begin by asking what problem the idea solves, what information enters the process, what result should leave it, and what evidence would show that the result is useful. In Transformers, LLMs, and Retrieval, the terms Evaluation are not vocabulary to memorize in isolation; they describe a system of relationships. A strong learner can translate the idea between plain language, a diagram, a small numerical or code example, and a statement of limitations. This page establishes that translation before adding detail.

03 REASONING

Central question: When should evaluation be used, and what observable evidence would make that choice defensible? Build a small local example that applies evaluation, records the result, and tests one failure condition.

04 IMPLEMENTATION

```
for batch in train_loader:
    optimizer.zero_grad()
    prediction = model(batch.inputs)
    loss = objective(prediction, batch.targets)
    loss.backward()
    optimizer.step()
# validate on held-out data after the epoch
```

05 PRACTICE

1. Define Evaluation in one precise sentence and name one assumption.
2. Create a two-step example of evaluation and verify the intermediate state.
3. Name one boundary case, one likely misconception, and one test that would expose it.
4. Connect this page to a prior concept and explain why the connection matters.



UNIT 08 / WORKFLOW

Evaluation: Workflow

01 OBJECTIVE

Explain and apply evaluation using explicit inputs, assumptions, steps, evidence, and failure checks.

02 CORE EXPLANATION

Evaluation is a central part of transformers, llms, and retrieval because it changes how inputs, assumptions, implementation choices, and evidence connect. A first-principles view separates the mechanism into state, operation, constraint, and test. State describes what is currently known. The operation transforms that state. A constraint rules out invalid moves, and a test compares the output with an expectation. This structure prevents an implementation from becoming a collection of unexplained steps. It also makes debugging local: identify which state, operation, constraint, or test failed. Apply this model directly to evaluation.

03 REASONING

Workflow: define the smallest valid input; predict the next state; perform one operation; test one invariant; then expand. Example: Build a small local example that applies evaluation, records the result, and tests one failure condition.

04 IMPLEMENTATION

```
for batch in train_loader:
    optimizer.zero_grad()
    prediction = model(batch.inputs)
    loss = objective(prediction, batch.targets)
    loss.backward()
    optimizer.step()
# validate on held-out data after the epoch
```

05 PRACTICE

1. Define Evaluation in one precise sentence and name one assumption.
2. Create a two-step example of evaluation and verify the intermediate state.
3. Name one boundary case, one likely misconception, and one test that would expose it.
4. Connect this page to a prior concept and explain why the connection matters.



UNIT 08 / WORKED EXAMPLE

Evaluation: Worked Example

01 OBJECTIVE

Explain and apply evaluation using explicit inputs, assumptions, steps, evidence, and failure checks.

02 CORE EXPLANATION

Consider this task: Build a small local example that applies evaluation, records the result, and tests one failure condition. First state the expected output before calculating or coding. Next choose a minimal representation and execute one step at a time. After every step, compare the intermediate state with the rule that should still hold. Finally, test a boundary case that differs from the example in exactly one important way. This worked-example pattern turns evaluation into a repeatable method rather than a memorized answer.

03 REASONING

Worked reasoning: 1) restate the task; 2) identify knowns and unknowns; 3) choose the smallest method; 4) calculate or trace; 5) verify independently; 6) explain what would make the result fail.

04 IMPLEMENTATION

```
for batch in train_loader:
    optimizer.zero_grad()
    prediction = model(batch.inputs)
    loss = objective(prediction, batch.targets)
    loss.backward()
    optimizer.step()
# validate on held-out data after the epoch
```

05 PRACTICE

1. Define Evaluation in one precise sentence and name one assumption.
2. Create a two-step example of evaluation and verify the intermediate state.
3. Name one boundary case, one likely misconception, and one test that would expose it.
4. Connect this page to a prior concept and explain why the connection matters.



UNIT 08 / FAILURE ANALYSIS

Evaluation: Failure Analysis

01 OBJECTIVE

Explain and apply evaluation using explicit inputs, assumptions, steps, evidence, and failure checks.

02 CORE EXPLANATION

Failure analysis asks how evaluation can produce a plausible but wrong result. Common causes include an incorrect assumption, a representation that loses information, an edge case omitted from tests, an invalid comparison, or evidence taken from the wrong stage of the process. Change one factor at a time, preserve the failing example, and write a regression test before repairing the implementation. A clear failure record is part of the learning artifact.

03 REASONING

Failure probe: construct the smallest input that breaks the naive approach to evaluation. State the expected behavior and why the naive result is misleading.

04 IMPLEMENTATION

```
for batch in train_loader:
    optimizer.zero_grad()
    prediction = model(batch.inputs)
    loss = objective(prediction, batch.targets)
    loss.backward()
    optimizer.step()
# validate on held-out data after the epoch
```

05 PRACTICE

1. Define Evaluation in one precise sentence and name one assumption.
2. Create a two-step example of evaluation and verify the intermediate state.
3. Name one boundary case, one likely misconception, and one test that would expose it.
4. Connect this page to a prior concept and explain why the connection matters.



UNIT 08 / APPLIED LAB

Evaluation: Applied Lab

01 OBJECTIVE

Explain and apply evaluation using explicit inputs, assumptions, steps, evidence, and failure checks.

02 CORE EXPLANATION

Implementation converts the idea of evaluation into inspectable behavior. Start with a deterministic example small enough to check by hand. Keep data preparation, core logic, and reporting separate. Add assertions for shape, type, range, or state as appropriate. Record the configuration and avoid relying on hidden global state. The goal is not merely to obtain output; it is to create an output whose origin can be explained and reproduced.

03 REASONING

Lab brief: Build a small local example that applies evaluation, records the result, and tests one failure condition. Save the input, expected output, observed output, and a short explanation of any difference.

04 IMPLEMENTATION

```
for batch in train_loader:
    optimizer.zero_grad()
    prediction = model(batch.inputs)
    loss = objective(prediction, batch.targets)
    loss.backward()
    optimizer.step()
# validate on held-out data after the epoch
```

05 PRACTICE

1. Define Evaluation in one precise sentence and name one assumption.
2. Create a two-step example of evaluation and verify the intermediate state.
3. Name one boundary case, one likely misconception, and one test that would expose it.
4. Connect this page to a prior concept and explain why the connection matters.



UNIT 09 / CONCEPT

Hallucination Controls: Concept

01 OBJECTIVE

Explain and apply hallucination controls using explicit inputs, assumptions, steps, evidence, and failure checks.

02 CORE EXPLANATION

Hallucination Controls is a central part of transformers, llms, and retrieval because it changes how inputs, assumptions, implementation choices, and evidence connect. Begin by asking what problem the idea solves, what information enters the process, what result should leave it, and what evidence would show that the result is useful. In Transformers, LLMs, and Retrieval, the terms Hallucination, Controls are not vocabulary to memorize in isolation; they describe a system of relationships. A strong learner can translate the idea between plain language, a diagram, a small numerical or code example, and a statement of limitations. This page establishes that translation before adding detail.

03 REASONING

Central question: When should hallucination controls be used, and what observable evidence would make that choice defensible? Build a small local example that applies hallucination controls, records the result, and tests one failure condition.

04 IMPLEMENTATION

```
for batch in train_loader:
    optimizer.zero_grad()
    prediction = model(batch.inputs)
    loss = objective(prediction, batch.targets)
    loss.backward()
    optimizer.step()
# validate on held-out data after the epoch
```

05 PRACTICE

1. Define Hallucination Controls in one precise sentence and name one assumption.
2. Create a two-step example of hallucination controls and verify the intermediate state.
3. Name one boundary case, one likely misconception, and one test that would expose it.
4. Connect this page to a prior concept and explain why the connection matters.



UNIT 09 / WORKFLOW

Hallucination Controls: Workflow

01 OBJECTIVE

Explain and apply hallucination controls using explicit inputs, assumptions, steps, evidence, and failure checks.

02 CORE EXPLANATION

Hallucination Controls is a central part of transformers, llms, and retrieval because it changes how inputs, assumptions, implementation choices, and evidence connect. A first-principles view separates the mechanism into state, operation, constraint, and test. State describes what is currently known. The operation transforms that state. A constraint rules out invalid moves, and a test compares the output with an expectation. This structure prevents an implementation from becoming a collection of unexplained steps. It also makes debugging local: identify which state, operation, constraint, or test failed. Apply this model directly to hallucination controls.

03 REASONING

Workflow: define the smallest valid input; predict the next state; perform one operation; test one invariant; then expand. Example: Build a small local example that applies hallucination controls, records the result, and tests one failure condition.

04 IMPLEMENTATION

```
for batch in train_loader:
    optimizer.zero_grad()
    prediction = model(batch.inputs)
    loss = objective(prediction, batch.targets)
    loss.backward()
    optimizer.step()
# validate on held-out data after the epoch
```

05 PRACTICE

1. Define Hallucination Controls in one precise sentence and name one assumption.
2. Create a two-step example of hallucination controls and verify the intermediate state.
3. Name one boundary case, one likely misconception, and one test that would expose it.
4. Connect this page to a prior concept and explain why the connection matters.



UNIT 09 / WORKED EXAMPLE

Hallucination Controls: Worked Example

01 OBJECTIVE

Explain and apply hallucination controls using explicit inputs, assumptions, steps, evidence, and failure checks.

02 CORE EXPLANATION

Consider this task: Build a small local example that applies hallucination controls, records the result, and tests one failure condition. First state the expected output before calculating or coding. Next choose a minimal representation and execute one step at a time. After every step, compare the intermediate state with the rule that should still hold. Finally, test a boundary case that differs from the example in exactly one important way. This worked-example pattern turns hallucination controls into a repeatable method rather than a memorized answer.

03 REASONING

Worked reasoning: 1) restate the task; 2) identify knowns and unknowns; 3) choose the smallest method; 4) calculate or trace; 5) verify independently; 6) explain what would make the result fail.

04 IMPLEMENTATION

```
for batch in train_loader:
    optimizer.zero_grad()
    prediction = model(batch.inputs)
    loss = objective(prediction, batch.targets)
    loss.backward()
    optimizer.step()
# validate on held-out data after the epoch
```

05 PRACTICE

1. Define Hallucination Controls in one precise sentence and name one assumption.
2. Create a two-step example of hallucination controls and verify the intermediate state.
3. Name one boundary case, one likely misconception, and one test that would expose it.
4. Connect this page to a prior concept and explain why the connection matters.



UNIT 09 / FAILURE ANALYSIS

Hallucination Controls: Failure Analysis

01 OBJECTIVE

Explain and apply hallucination controls using explicit inputs, assumptions, steps, evidence, and failure checks.

02 CORE EXPLANATION

Failure analysis asks how hallucination controls can produce a plausible but wrong result. Common causes include an incorrect assumption, a representation that loses information, an edge case omitted from tests, an invalid comparison, or evidence taken from the wrong stage of the process. Change one factor at a time, preserve the failing example, and write a regression test before repairing the implementation. A clear failure record is part of the learning artifact.

03 REASONING

Failure probe: construct the smallest input that breaks the naive approach to hallucination controls. State the expected behavior and why the naive result is misleading.

04 IMPLEMENTATION

```
for batch in train_loader:
    optimizer.zero_grad()
    prediction = model(batch.inputs)
    loss = objective(prediction, batch.targets)
    loss.backward()
    optimizer.step()
# validate on held-out data after the epoch
```

05 PRACTICE

1. Define Hallucination Controls in one precise sentence and name one assumption.
2. Create a two-step example of hallucination controls and verify the intermediate state.
3. Name one boundary case, one likely misconception, and one test that would expose it.
4. Connect this page to a prior concept and explain why the connection matters.



UNIT 09 / APPLIED LAB

Hallucination Controls: Applied Lab

01 OBJECTIVE

Explain and apply hallucination controls using explicit inputs, assumptions, steps, evidence, and failure checks.

02 CORE EXPLANATION

Implementation converts the idea of hallucination controls into inspectable behavior. Start with a deterministic example small enough to check by hand. Keep data preparation, core logic, and reporting separate. Add assertions for shape, type, range, or state as appropriate. Record the configuration and avoid relying on hidden global state. The goal is not merely to obtain output; it is to create an output whose origin can be explained and reproduced.

03 REASONING

Lab brief: Build a small local example that applies hallucination controls, records the result, and tests one failure condition. Save the input, expected output, observed output, and a short explanation of any difference.

04 IMPLEMENTATION

```
for batch in train_loader:
    optimizer.zero_grad()
    prediction = model(batch.inputs)
    loss = objective(prediction, batch.targets)
    loss.backward()
    optimizer.step()
# validate on held-out data after the epoch
```

05 PRACTICE

1. Define Hallucination Controls in one precise sentence and name one assumption.
2. Create a two-step example of hallucination controls and verify the intermediate state.
3. Name one boundary case, one likely misconception, and one test that would expose it.
4. Connect this page to a prior concept and explain why the connection matters.



FILE / ORIENTATION

Deep-Dive Capstone

01 FILE GUIDANCE

Build a compact, reproducible artifact demonstrating transformers, llms, and retrieval. Include assumptions, a baseline, tests, failure analysis, results, limitations, and instructions for repeating the work offline.

02 LOCAL USE

Index this page in the offline database and preserve its resource and page identifiers for bookmarks, notes, highlights, and progress.